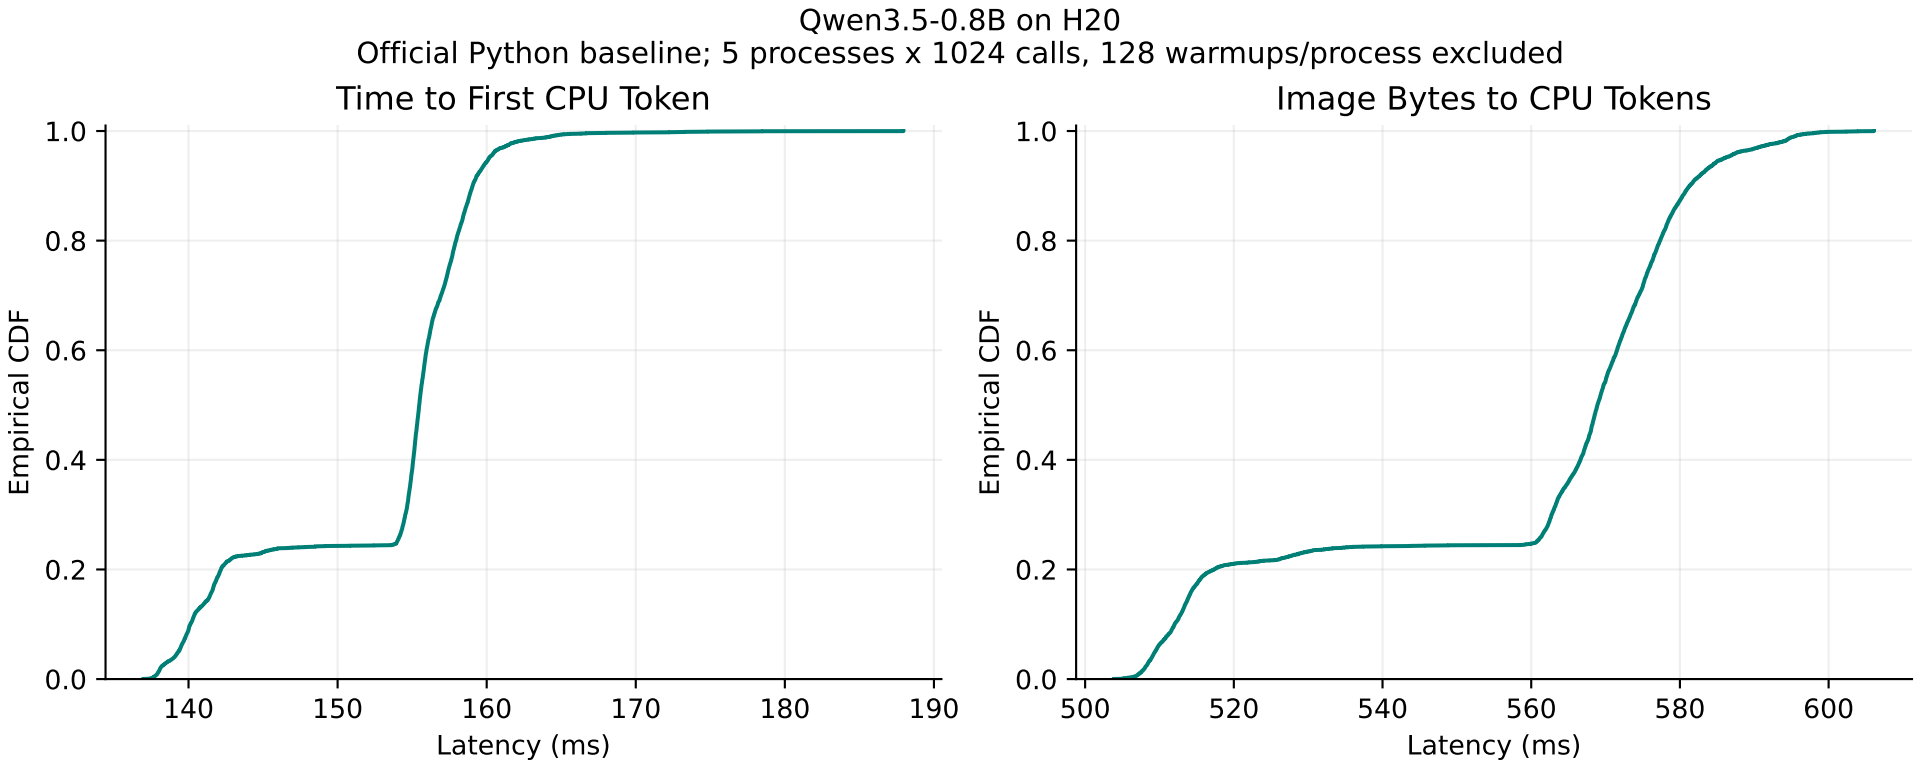
Qwen3.5-0.8B on H20
Official Python baseline; 5 processes x 1024 calls, 128 warmups/process excluded
Time to First CPU Token
Image Bytes to CPU Tokens
Empirical CDF
Latency (ms)
0.0
0.2
0.4
0.6
0.8
1.0
140
150
160
170
180
190
500
520
540
560
580
600